



BITS Pilani
Pilani Campus

Network Programming

K Hari Babu
Department of Computer Science & Information Systems

Outline



- Sockets
- TCP Client Server




BITS Pilani
Pilani Campus



Sockets

OSI & Internet protocol suite



TCP/IP & Sockets API



- TCP/IP does not include an API definition.
- There are a variety of APIs for use with TCP/IP:
 - Sockets
 - TLI, XTI
 - Winsock
 - MacTCP
- API should have the following functionalities
 - Specify local and remote communication endpoints
 - Initiate a connection
 - Wait for incoming connection
 - Send and receive data
 - Terminate a connection gracefully
 - Error handling

Berkeley Sockets API



- First appeared in 42. BSD in 1983.
 - Supported on every UNIX variant.
 - WinSock API also follows socket API.
- Generic API:
 - support for multiple communication domains which differ in addressing methods.

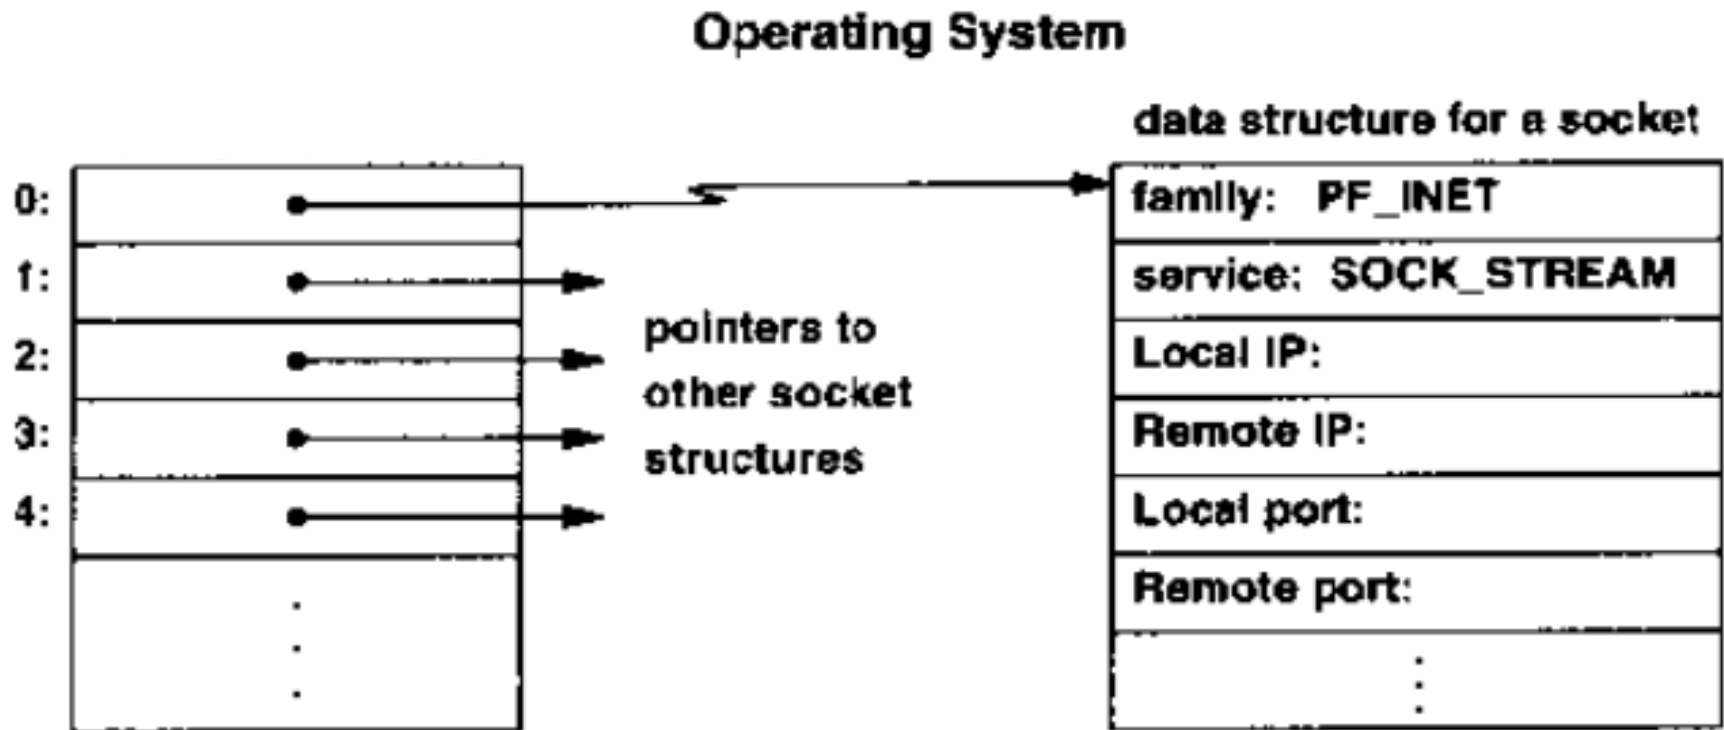
Domain	Communication performed	Communication between applications	Address format	Address structure
AF_UNIX	within kernel	on same host	pathname	<i>sockaddr_un</i>
AF_INET	via IPv4	on hosts connected via an IPv4 network	32-bit IPv4 address + 16-bit port number	<i>sockaddr_in</i>
AF_INET6	via IPv6	on hosts connected via an IPv6 network	128-bit IPv6 address + 16-bit port number	<i>sockaddr_in6</i>

- Uses existing I/O programming interface as much as possible.
 - Socket API is similar to file I/O

- A socket is an abstract representation of a communication endpoint.
- Sockets work with Unix I/O services just like files, pipes & FIFOs.
 - read(), write() and close() sys calls work on sockets also
- Sockets have special needs over files:
 - establishing a connection.
 - specifying communication endpoint addresses.

- Two types of sockets
 - Stream Sockets
 - Reliable
 - Bidirectional
 - Byte-stream
 - Connection-oriented
 - Stream sockets operate in connected pairs. (local end point, remote end point)
 - Internet Domain: TCP Socket
 - Datagram Sockets
 - Message boundaries are preserved.
 - No reliability support: out of order, duplicates, or lost datagrams
 - Connectionless socket: no need to be connected to another socket.
 - Internet Domain: UDP Socket

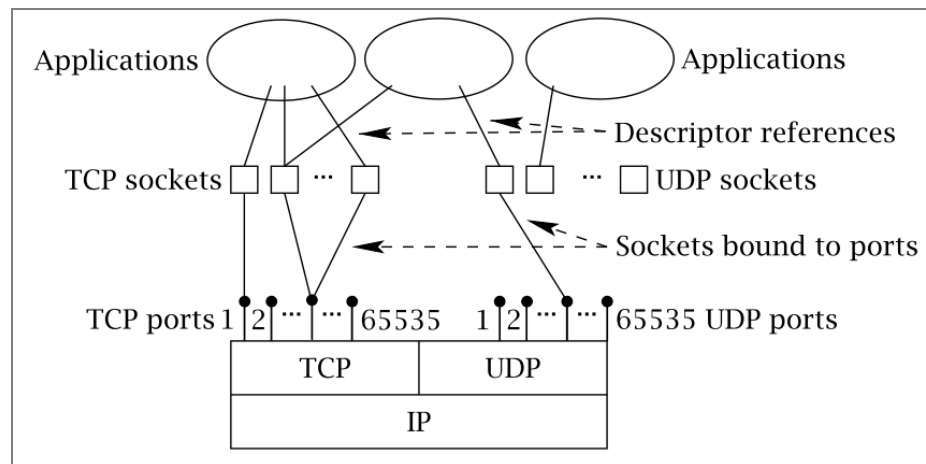
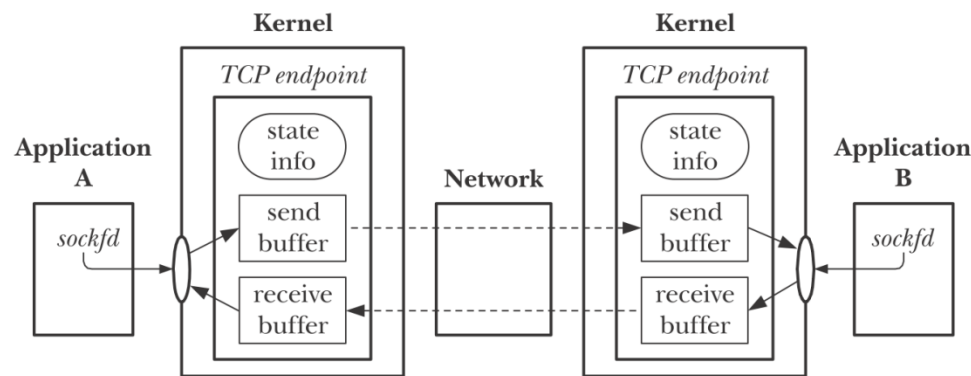
Socket Descriptor Data Structure



Internet Domain & Sockets




- Endpoint is identified by ip address and port number.
 - A socket needs to be bound to endpoints local and remote.



Socket Pair




- The socket pair for a TCP connection is the four-tuple that defines the two endpoints of the connection:
 - the local IP address, local port, foreign IP address, and foreign port.
- A socket pair uniquely identifies every TCP connection on a network.
- We can extend the concept of a socket pair to UDP, even though UDP is connectionless.
- TCP connection for a ftp server:
 - Server has two IP interfaces. * indicates any ip address/any port.




TCP Connection



- A client 206.168.112.219 connects to 12.106.32.254 at port 21.



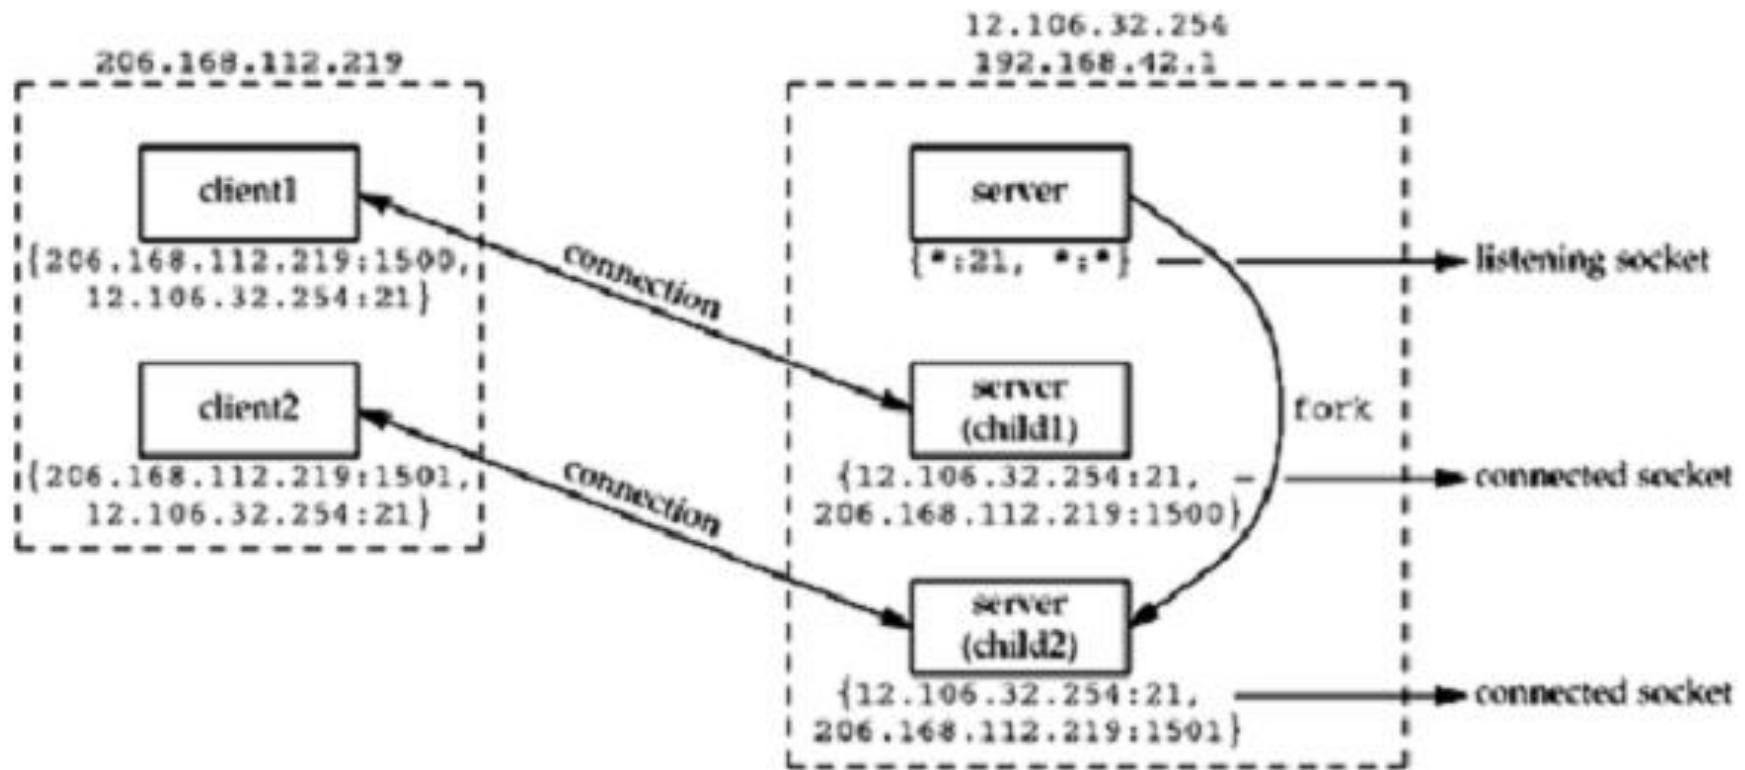
- Server creates a child to handle new client.
 - See the new TCP connection pair on server.




TCP Connection



- Client opens one more connection to the same server. What makes it a different TCP connection?



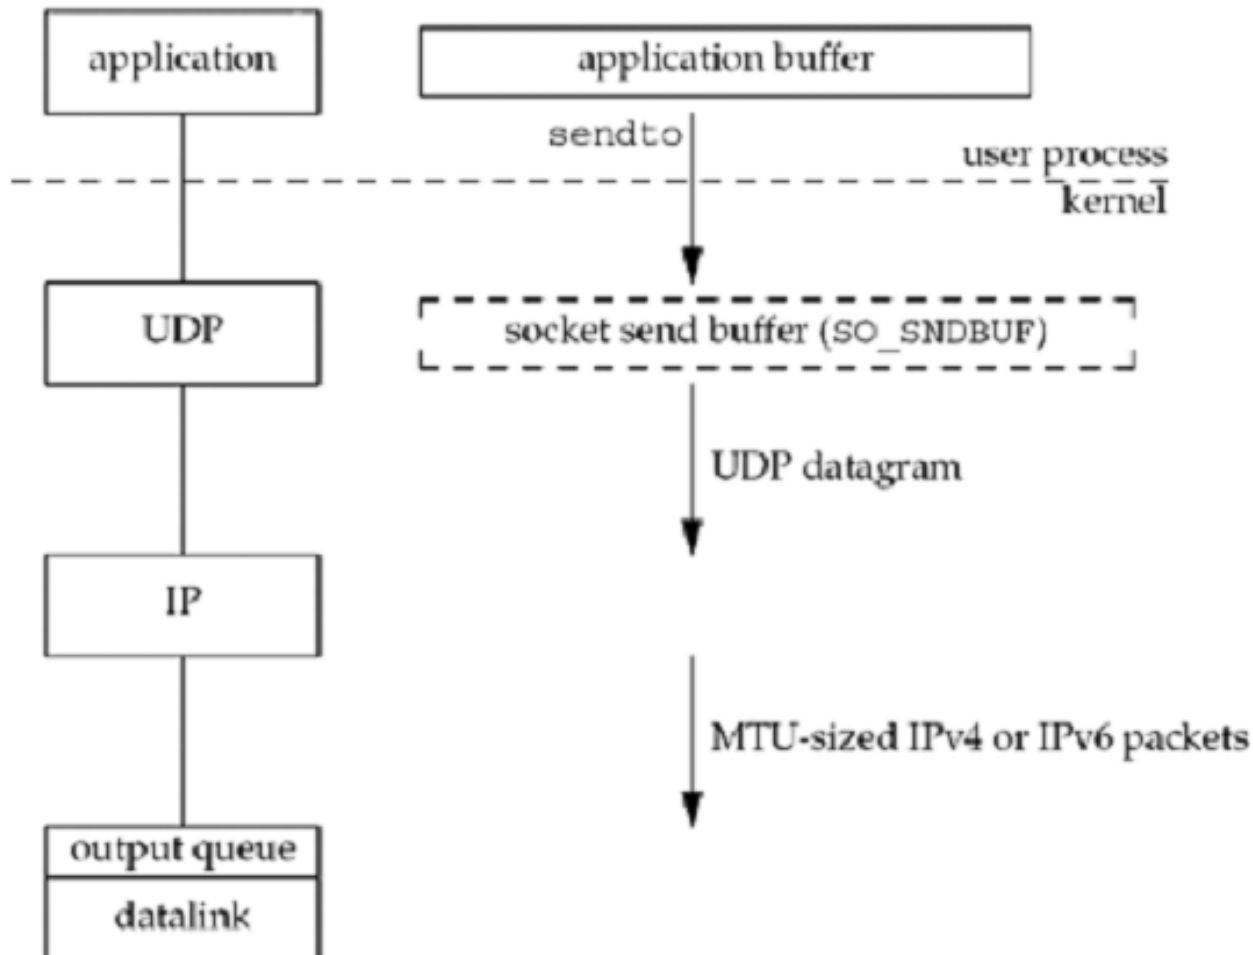
Writing to TCP Socket



Writing to UDP Socket



UDP socket.



Creating a Socket



```
1  #include <sys/socket.h>
2  int socket(int domain , int type , int protocol );
3  //Returns file descriptor on success, or -1 on error
```

- Domain argument specifies the communication domain.
- Type argument specifies the socket type
- Protocol is usually 0. Its value is automatically from first two arguments.
 - But for raw sockets, protocol value can be specified.
- The **socket()** system call returns a socket descriptor (small integer) or -1 on error.
- **socket()** allocates resources needed for a communication endpoint - but it does not deal with endpoint addressing.

Specifying an Endpoint Address



- Remember that the sockets API is generic.
- There must be a generic way to specify endpoint addresses.
- Internet Domain requires an IP address and a port number for each endpoint address.
- Other domains (families) may use other schemes.
- Data types for specifying address structure

```
1  sa_family_t  //address family
2  socklen_t    //length of struct
3  in_addr_t    //IPv4 address
4  in_port_t    //IP port number
```

Generic Socket Address Structure



- General socket address structure acts as template.
- All sys calls take this as the input.

```
1 struct sockaddr {  
2     sa_family_t sa_family; /* Address family (AF_* constant) */  
3     char        sa_data[14]; /* Socket address (size varies  
4                             according to socket domain) */  
5 };
```

- For IPv4 domain, socket address structure is

```
1 struct sockaddr_in {  
2     sa_family_t    sin_family;  
3     in_port_t      sin_port;  
4     struct in_addr sin_addr;  
5     char           sin_zero[8];  
6 };
```

```
1 sa_family_t //address family  
2 socklen_t   //length of struct  
3 in_addr_t   //IPv4 address  
4 in_port_t   //IP port number
```

```
2 struct in_addr {  
3     in_addr_t s_addr;  
4 };  
5
```

Specifying End Point Address



- Bind call binds an address to a socket

```
1 #include <sys/socket.h>
2 int bind(int sockfd , const struct sockaddr * addr ,
3 socklen_t addrlen );
4 //Returns 0 on success, or -1 on error
```

- Bind is used to bind the socket to a local end point.
 - *addr* is the pointer to an address structure that contains the endpoint address.


process specifies		result
IP address	port	
Wildcard	0	kernel chooses IP addr and port
Wild card	nonzero	kernel chooses IP, process specifies port
local IP addr	0	process specifies IP, kernel chooses port
local IP addr	nonzero	process specifies IP and port

Wildcard specified as INADDR_ANY

Byte Order



- Different hardware architectures in the network. Passing integers across the network needs to be in standard format: Network Byte Order (NBO): bigendian byte order.



MSB = Most Significant Byte, LSB = Least Significant Byte

Network Byte Order Functions



```
1  #include <arpa/inet.h>
2  uint16_t htons(uint16_t  host_uint16 );
3  //Returns host_uint16 converted to network byte order
4  uint32_t htonl(uint32_t  host_uint32 );
5  //Returns host_uint32 converted to network byte order
6  /uint16_t ntohs(uint16_t  net_uint16 );
7  //Returns net_uint16 converted to host byte order
8  uint32_t ntohl(uint32_t  net_uint32 );
9  //Returns net_uint32 converted to host byte order
```

‘h’ : host byte order

‘s’ : short (16bit)

‘n’ : network byte order

‘l’ : long (32bit)

bind() Example



```
1  int mysock, err;  
2  struct sockaddr_in myaddr;  
3  
4  mysock = socket(PF_INET, SOCK_STREAM, 0);  
5  myaddr.sin_family = AF_INET;  
6  myaddr.sin_port = htons( portnum );  
7  myaddr.sin_addr = htonl( ipaddress );  
8  
9  err = bind(mysock, (sockaddr *) &myaddr, sizeof(myaddr));
```

- Note the usage of byte conversion.
- Note the casting internet domain address to generic address structure.
- There are a number of uses for **bind()**:
 - Server would like to bind to a well known address (port number).
 - Client can bind to a specific port.
 - Client can ask the O.S. to assign *any available* port number.

IPv4 Address Conversion



- The `inet_pton()` and `inet_ntop()` functions allow conversion of both IPv4 and IPv6 addresses between binary form and dotted-decimal or hex-string notation.
- The `p` in the names of these functions stands for “presentation,” and the `n` stands for “network.” The presentation form is a human-readable string.

```
1  #include <arpa/inet.h>
2  int inet_pton(int domain , const char * src_str , void * addrptr );
3  //Returns 1 on successful conversion, 0 if src_str is not in
4  //presentation format, or -1 on error
5  const char *inet_ntop(int domain , const void * addrptr ,
6  char * dst_str , size_t len );
7  //Returns pointer to dst_str on success, or NULL on error
```



BITS Pilani
Pilani Campus



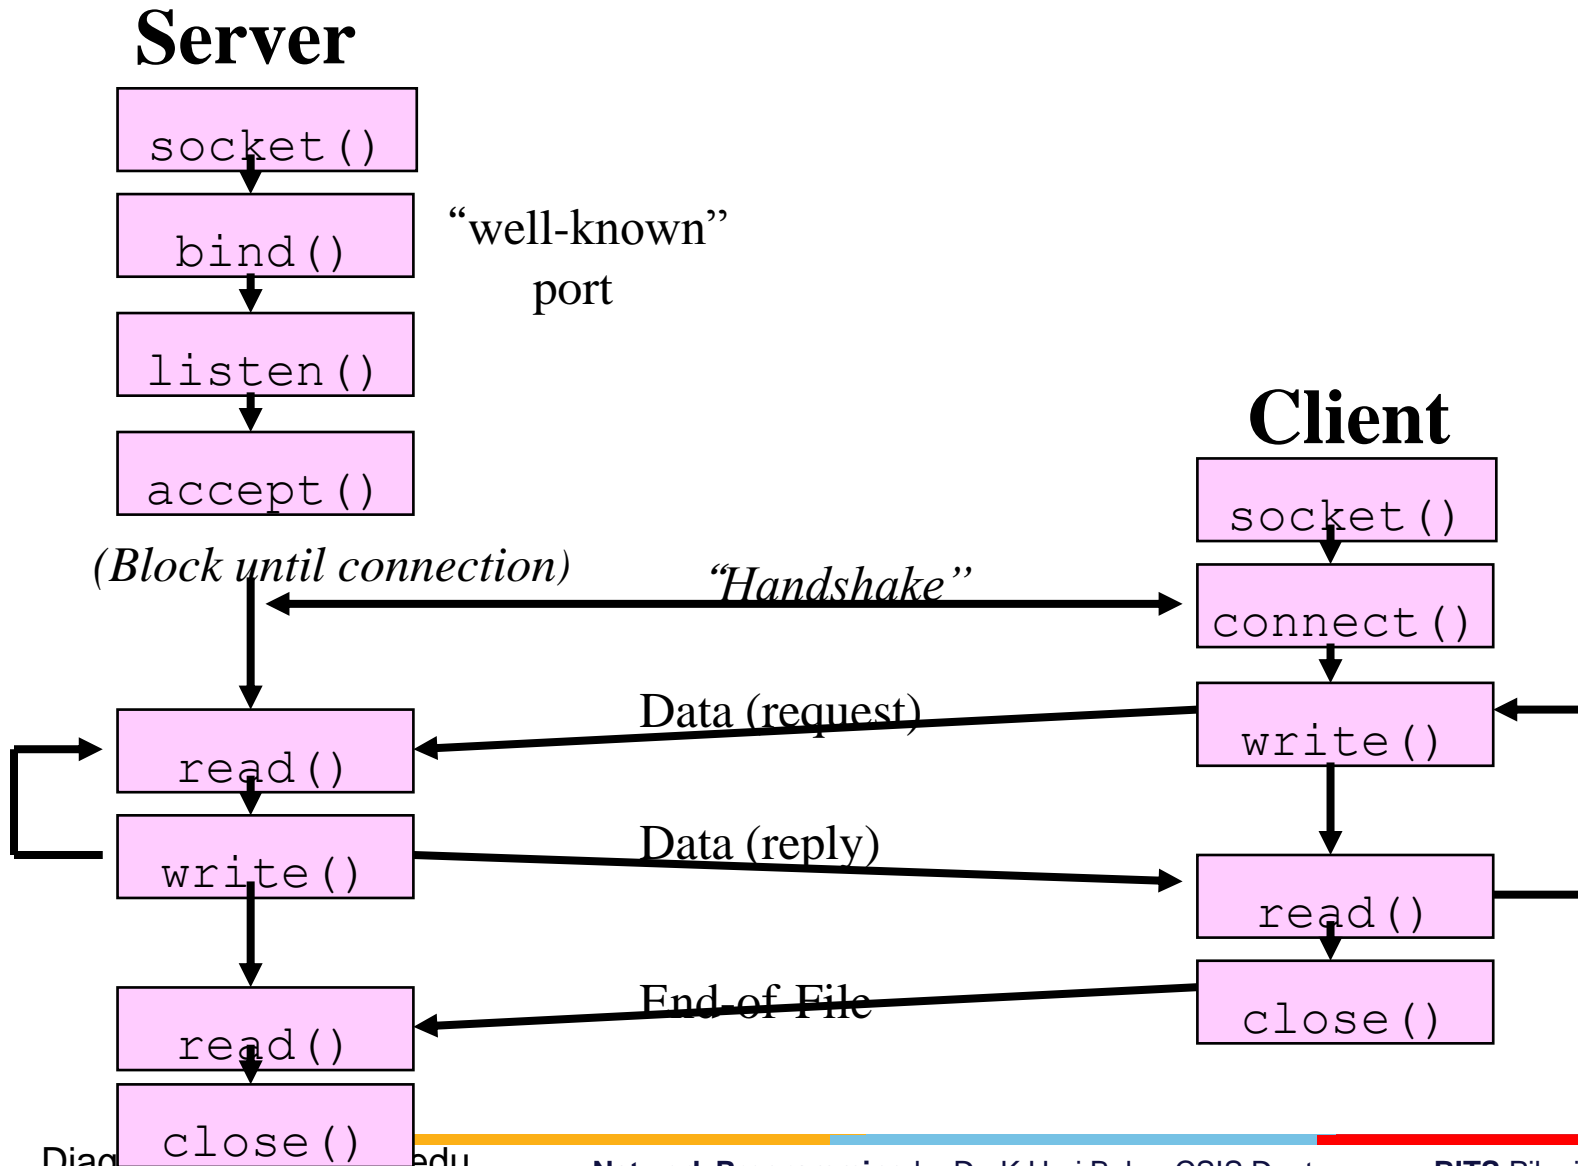
Client Server Programming

Active & Passive Sockets

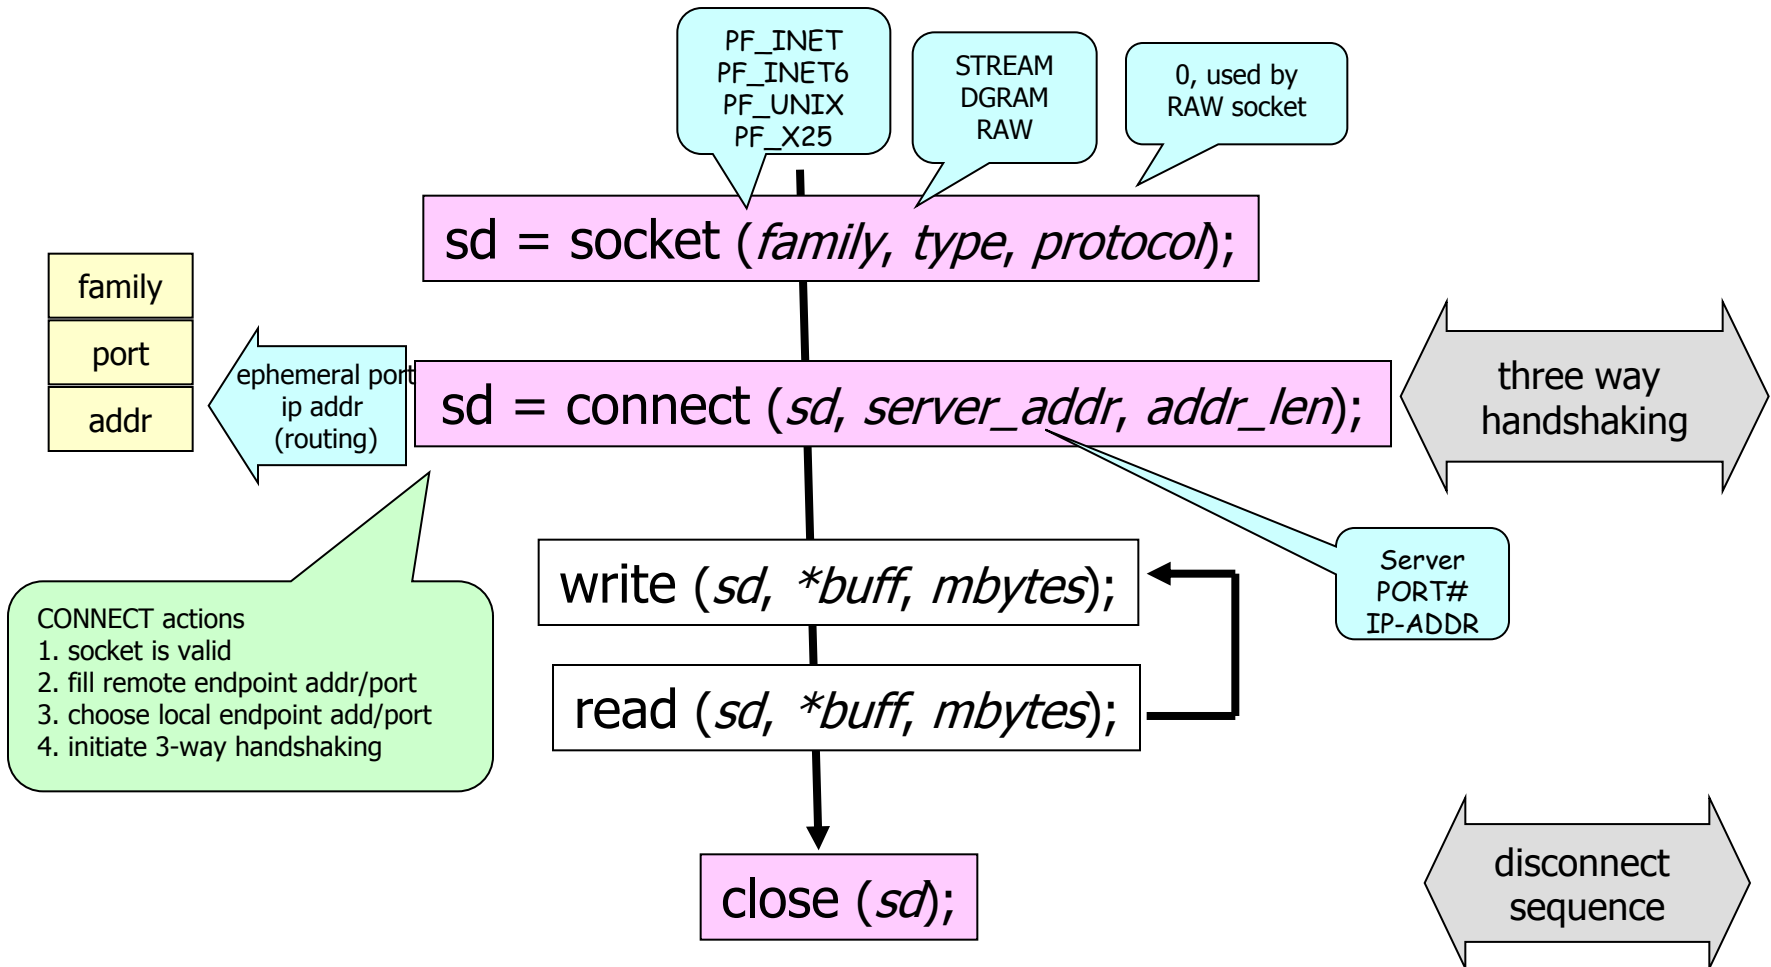


- By default, a socket that has been created using `socket()` is active.
 - An active socket can be used in a `connect()` call to establish a connection to a passive socket.
 - This is referred to as performing an active open.
- A passive socket (also called a listening socket) is one that has been marked to allow incoming connections by calling `listen()`.
 - Accepting an incoming connection is referred to as performing a passive open.
- client:
 - which does active socket open
- Server:
 - which does passive socket open.

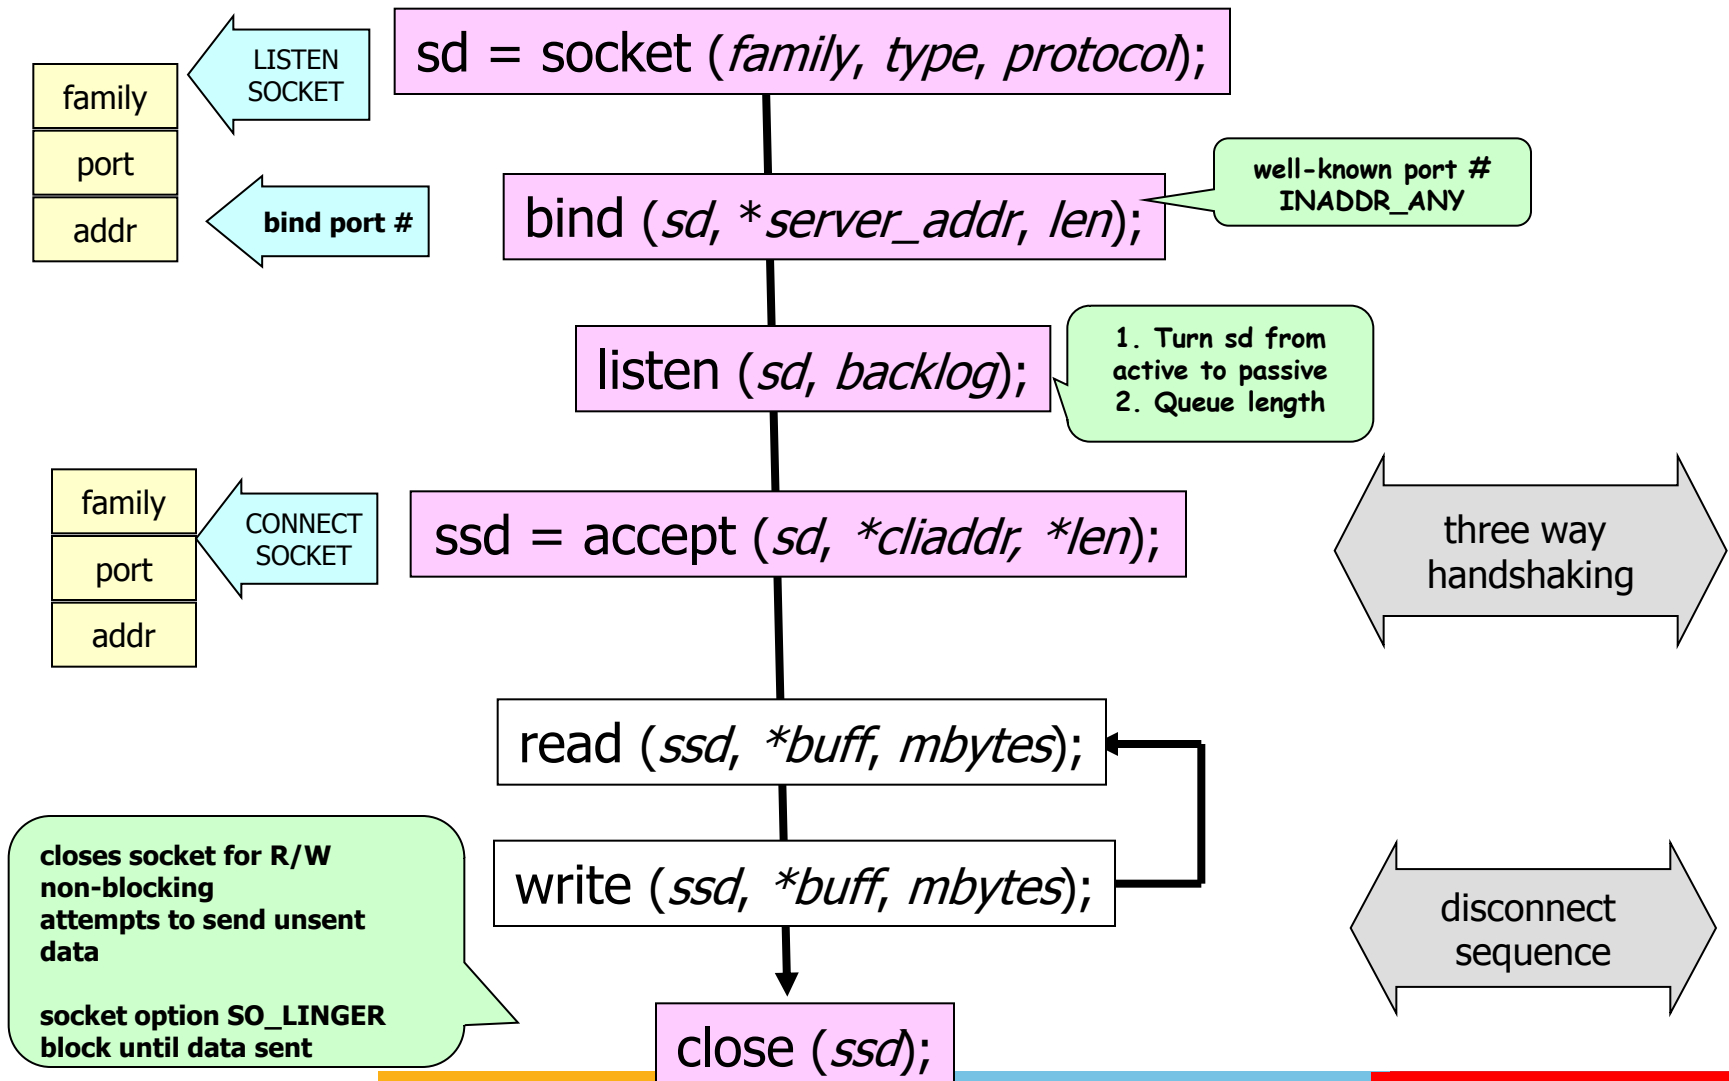
TCP Client Server



TCP Client



TCP Server



connect() - connect to server



- The connect() system call connects the active socket referred to by the file descriptor sockfd to the listening socket whose address is specified by addr and addrlen.

```
1  #include <sys/socket.h>
2  int connect(int sockfd , const struct sockaddr * addr,
3  socklen_t addrlen );
4  //Returns 0 on success, or -1 on error
```

- *sockfd* is socket descriptor from `socket()`
- *servaddr* is a pointer to a structure with:
 - *port number* and *IP address*
 - must be specified (unlike `bind()`)
- *addrlen* is length of structure
- client doesn't need `bind()`
 - OS will pick ephemeral port
- returns socket descriptor if ok, -1 on error

connect() - connect to server



- Errors

- If the server's TCP response to client TCP's SYN segment is RST, then there is no process is waiting for incoming connections.
 - Hard error
- Three conditions that generate RST are
 - when a SYN arrives for a port that has no listening server
 - when TCP wants to abort an existing connection
 - when TCP receives a segment for a connection that does not exist
- If client's SYN request elicits ICMP "destination unreachable" message, kernal saves the message but keeps on sending the SYN segment.
 - Soft error

listen() - change socket state to passive





- The listen() system call marks the stream socket referred to by the file descriptor sockfd as passive.
 - The socket will subsequently be used to accept connections from other (active) sockets.

```
1  #include <sys/socket.h>
2  int listen(int  sockfd , int  backlog );
3  //Returns 0 on success, or -1 on error
```

- *sockfd* is socket descriptor from `socket()`
- *backlog* is maximum number of connections that the server should queue for this socket
 - historically 5
 - rarely above 15 on a even moderate Web server!

listen()



accept() - return next completed connection



- The `accept()` system call accepts an incoming connection on the listening stream socket referred to by the file descriptor `sockfd`.
 - If there are no pending connections when `accept()` is called, the call blocks until a connection request arrives.

```
1  #include <sys/socket.h>
2  int accept(int  sockfd , struct sockaddr * addr,
3  socklen_t * addrlen );
4  //Returns file descriptor on success, or -1 on error
```

- *sockfd* is socket descriptor from `socket()`
- *cliaddr* and *addrlen* return protocol address from client
- returns brand new descriptor, created by OS
 - if used with `fork()` , can create concurrent server

close() - close socket fd



- *sockfd* is socket descriptor from `socket()`
- closes socket for reading/writing
 - returns (doesn't block)
 - attempts to send any unsent data
 - socket option `SO_LINGER`
 - block until data sent
 - or discard any remaining data
 - Returns -1 if error

Descriptor Reference Counts



- For every socket a reference count is maintained, as to how many processes are accessing that socket
- When `close()` is called on socket descriptor reference count is decreased by 1
- When `close()` is called on socket descriptor, TCP 4 packet termination sequence will be initiated only if the reference count goes to zero.

getsockname() and getpeername() Functions



- getsockname return the local endpoint address associated with a socket
- getpeername return the foreign protocol address associated with a socket

```
1  #include <sys/socket.h>
2  int getsockname(int sockfd, struct sockaddr *localaddr,
3  socklen_t *addrlen);
4  int getpeername(int sockfd, struct sockaddr *peeraddr,
5  socklen_t *addrlen);
```

TCP Echo Client




- TCP client and server using echo protocol

```
1  int main(int argc, char **argv)
2  {
3      int      sockfd;
4      struct sockaddr_in servaddr;
5      if (argc != 2)
6          err_quit("usage: tcpcli <IPaddress>");
7      sockfd = Socket(PF_INET, SOCK_STREAM, 0);
8      bzero(&servaddr, sizeof(servaddr));
9      servaddr.sin_family = AF_INET;
10     servaddr.sin_port = htons(SERV_PORT);
11     inet_pton(AF_INET, argv[1], &servaddr.sin_addr);
12     connect(sockfd, (SA *)&servaddr, sizeof(servaddr));
13     str_cli(stdin, sockfd);

17 void str_cli(FILE *fp, int sockfd)
18 {
19     char sendline[MAXLINE], recvline[MAXLINE];
20     while (Fgets(sendline, MAXLINE, fp) != NULL) {
21         write(sockfd, sendline, strlen(sendline));
22         if (wead(sockfd, recvline, MAXLINE) == 0)
23             err_quit("str_cli: server terminated prematurely");
24         wputs(recvline, stdout);
25     }
```

TCP Concurrent Server



TCP Concurrent Server



```
1  int main(int argc, char **argv)
2  {  int      listenfd, connfd;
3     pid_t    childpid;
4     socklen_t clilen;
5     struct sockaddr_in cliaddr, servaddr;
6     listenfd = Socket (AF_INET, SOCK_STREAM, 0);
7     bzero(&servaddr, sizeof(servaddr));
8     servaddr.sin_family = AF_INET;
9     servaddr.sin_addr.s_addr = htonl (INADDR_ANY);
10    servaddr.sin_port = htons (SERV_PORT);
11    bind(listenfd, (SA *) &servaddr, sizeof(servaddr));
12    listen(listenfd, LISTENQ);
13    for ( ; ; ) {
14        clilen = sizeof(cliaddr);
15        connfd = accept(listenfd, (struct sockaddr *) &cliaddr, &clilen);
16        if ( (childpid = fork()) == 0) { /* child process */
17            close(listenfd); /* close listening socket */
18            str_echo(connfd); /* process the request */
19            exit (0);
20        }
21        close(connfd); /* parent closes connected socket */
22    }
23 }
```

str_echo function

innovate

achieve

lead

```
1 void str_echo(int sockfd)
2 {
3     ssize_t n;
4     char    buf[MAXLINE];
5     again:
6     while ( (n = read(sockfd, buf, MAXLINE)) > 0)
7         Write(sockfd, buf, n);
8     if (n < 0 && errno == EINTR)
9         goto again;
10    else if (n < 0)
11        err_sys("str_echo: read error");
12 }
```


TCP Concurrent Server



- Handling zombies
 - while ((pid = waitpid(-1, &stat, WNOHANG)) > 0) in SIGCHLD signal handler
- Handling interrupted system calls
 - when writing network programs that catch signals, we must be cognizant of interrupted system calls, and we must handle them
 - Slow system call is any system call that can block forever

Handling interrupted system calls



```
for ( ; ; ) {  
    clilen = sizeof (cliaddr);  
    if ( (connfd = accept (listenfd, (SA *)  
        &cliaddr, &clilen)) < 0) {  
        if (errno == EINTR)  
            continue;           /* back to for () */  
        else  
            err_sys ("accept error");  
    }  
}
```

- Another option:
 - Use sigaction() with SA_RESTART flag.

Termination of Server Process



- FIN is sent to client
- Client tcp sends ACK to server
- What if client application doesn't take note of it, and sends data to server?

SIGPIPE Signal



- When a process writes to a socket that has received an RST, the SIGPIPE signal is sent to the process.
- The default action of this signal is to terminate the process, so the process must catch the signal to avoid being involuntarily terminated.

Crashing of Server Host



- Nothing is sent to client
- Client will try to reach the host, but will get errors such as ETIMEDOUT, EHOSTUNREACH, ENETWORKUNREACH

Crashing and Rebooting of Server Host



- When client sends packets, server will respond with RST

Shutdown of Server Host



- Init sends SIGTERM to all processes
- Then sends SIG KILL to all processes
- Fin is sent to the client

Exercise



- Write a TCP client and server that fulfills the following requirements.
- Server.c:
 - server should take port number on command-line and listen on that port.
 - server should create a child to handle a new client.
 - When a client sends a command such as 'ps', server should execute the command and send the output to the client.
 - it should take care of zombies.
- Client.c:
 - client takes ip address and port number of the server on command-line.
 - client sets up a connection to the server.
 - client takes a command from the user and sends it to the server.
 - client waits for the reply and prints the reply on the standard output.


```
1  if ((servSock = socket(PF_INET, SOCK_STREAM, IPPROTO_TCP)) < 0)
2  printf("socket() failed");
3  memset(&echoServAddr, 0, sizeof(echoServAddr));
4  echoServAddr.sin_family = AF_INET;
5  echoServAddr.sin_addr.s_addr = htonl(INADDR_ANY);
6  echoServAddr.sin_port = htons(echoServPort);
7  if (bind
8  (servSock, (struct sockaddr *) &echoServAddr,
9   sizeof(echoServAddr)) < 0)
10 printf("bind() failed");
11 if (listen(servSock, MAXPENDING) < 0)
12 printf("listen() failed");
```



```
14  for (;;) {
15      clntLen = sizeof(echoClntAddr);
16      if ((clntSock = accept(servSock, (struct sockaddr *) &echoClntAddr,
17                          &clntLen)) < 0)//2M
18          printf ("accept() failed");
19      printf("Handling client %s\n", inet_ntoa(echoClntAddr.sin_addr));
20      ret=fork(); //3M
21      if(ret==0){
22          close(listenfd);
23          read(clntSock, buff, size);
24          int p[2];
25          pipe(p);
26          ret=fork();
27          if(ret==0)
28          {   close(1);
29              dup(p[1]);
30              execv(buff);
31          }
32          read(p[0],buff,size);
33          write(clntSock,buff,size);
34      }
35      close(clntSock);
36  }
```

```
1 //Step 1: Set up Address Structure
2 bzero(&Server_Address, sizeof(Server_Address));
3 Server_Address.sin_family = AF_INET;
4 Server_Address.sin_port = htons(port);
5 temp = inet_addr(Address);
6 if (temp != INADDR_NONE){
7     Server_Address.sin_addr.s_addr = temp;
8 }else{
9     printf ("Invalid IP Address.");
10 }
11 //Step 2: Create a Socket
12 mysocket = socket(PF_INET, SOCK_STREAM, 0);
13 if (mysocket == -1) printf ("socket()");
14
15 //Step 3: Connect to Server
16 result = connect(mysocket, (struct sockaddr *) &Server_Address,
17 sizeof(Server_Address) );
18 if (result == -1) printf ("connect()");
19 while(1){
20     scanf("%s", cmd);
21     if(strcmp(cmd,"exit")==0)
22         exit(0);
23     write(mysocket, cmd,(strlen(cmd)));
24     read(mysocket,buff,size);
25     printf("%s", buff);}}
```

Acknowledgements



Q&A





BITS Pilani
Pilani Campus



Thank You