



**BITS Pilani**  
Pilani Campus

# Network Programming

K Hari Babu  
Department of Computer Science & Information Systems



**BITS Pilani**  
Pilani Campus



# Outline

- I/O Multiplexing (cont..)
  - Client using select()
  - shutdown()
  - Concurrent server using select()
- Non-blocking sockets
  - Nonblocking read(), write()
  - Nonblocking connect()



# I/O Multiplexing

T1: Ch6

# I/O Multiplexing



- I/O multiplexing allows us to simultaneously monitor multiple file descriptors to see if I/O is possible on any of them.
- `select()`, appeared along with the sockets API in BSD. This was historically the more widespread of the two system calls. The other system call, `poll()`, appeared in System V.
- We can use `select()` and `poll()` to monitor file descriptors for regular files, terminals, pseudoterminals, pipes, FIFOs, sockets, and some types of character devices.
- Both system calls allow a process either to block indefinitely waiting for file descriptors to become ready or to specify a timeout on the call.

# select()



- The select() system call blocks until one or more of a set of file descriptors becomes ready.

```
1  #include <sys/time.h>    /* For portability */
2  #include <sys/select.h>
3  int select(int  nfds , fd_set * readfds , fd_set * writefds,
4             fd_set * exceptfds, struct timeval * timeout );
5  //Returns number of ready file descriptors, 0 on timeout, or -1 on error
```

- *nfds*: highest number assigned to a descriptor +1.
- *readfds*: set of descriptors we want to read from.
- *writefds*: set of descriptors we want to write to.
- *exceptfds*: set of descriptors to watch for exceptions.
- *timeout*: maximum time select should wait

```
7  struct timeval {
8      long tv_usec;    /* seconds */
9      long tv_usec;    /* microseconds */
10 }
```

# select()



- `timeval==NULL`
  - Wait forever : return only when descriptor is ready
- `timeval != NULL`: wait up to a fixed amount of time
  - `timeval = 0`
    - Do not wait at all : return immediately after checking the descriptors
  - `Timeval>0`
    - Return only if descriptor is ready or `timeval` expires.

# File descriptor sets



- The `readfds`, `writfds`, and `exceptfds` arguments are pointers to file descriptor sets, represented using the data type `fd_set`.
- the `fd_set` data type is implemented as a bit mask.

```
1  #include <sys/select.h>
2  void FD_ZERO(fd_set *fdset);
3  /* clear all bits in fdset */
4  void FD_SET(int fd, fd_set *fdset);
5  /* turn on the bit for fd in fdset */
6  void FD_CLR(int fd, fd_set *fdset);
7  /* turn off the bit for fd in fdset */
8  int FD_ISSET(int fd, fd_set *fdset);
9  /* is the bit for fd on in fdset ? */
10 //Returns true (1) if fd is in fdset, or false (0) otherwise
```

- A file descriptor set has a maximum size, defined by the constant `FD_SETSIZE`. On Linux, this constant has the value 1024.



# select()



- *nfds*
  - Its value is the maximum descriptor to be tested, plus one
    - example: fds 1,2,5 => nfds: 6
- *readset*
  - descriptor set for checking readable
- *writeset*
  - descriptor set for checking writable
- *exceptset*
  - descriptor set for checking two exception conditions
    - arrival of out of band data for a socket
    - the presence of control status information to be read from the master side of a pseudo terminal
- When select returns value > 1, these sets have been modified by kernel. Now they contain the fds which are ready.

# When is the descriptor ready for reading?



- The number of bytes of data in the socket receive buffer is greater than or equal to the current size of the low-water mark for the socket receive buffer. `SO_RCVLOWAT` socket option. It defaults to 1 for TCP and UDP sockets
- The read half of the connection is closed (i.e., a TCP connection that has received a FIN)
- The socket is a listening socket and the number of completed connections is nonzero.
- A socket error is pending. A read operation on the socket will not block and will return an error (`-1`) with `errno` set to the specific error condition.

# When the socket is ready for writing?




- The number of bytes of available space in the socket send buffer is greater than or equal to the current size of the low-water mark for the socket send buffer. 2048 bytes.
- The write half of the connection is closed. A write operation on the socket will generate SIGPIPE.
- A socket using a non-blocking connect has completed the connection, or the connect has failed
- A socket error is pending. A write operation on the socket will not block and will return an error (−1) with errno set to the specific error condition.
- These pending errors can also be fetched and cleared by calling getsockopt with the SO\_ERROR socket option.

# Client Handling Multiple Descriptors



- A client is handling two descriptors.
  - *stdin*
  - *socket*
- Sequential handling:
  - First wait on *stdin*.
  - Write to *socket*
  - Read from *socket*.
  - Write to *stdout*.




Normal course of actions

```
1 void str_cli(FILE *fp, int sockfd)
2 {
3     char sendline[MAXLINE], recvline[MAXLINE];
4     while (Fgets(sendline, MAXLINE, fp) != NULL) {
5         Writen(sockfd, sendline, strlen(sendline));
6         if (Readline(sockfd, recvline, MAXLINE) == 0)
7             err_quit("str_cli: server terminated prematurely");
8         Fputs(recvline, stdout);
9     }
10 }
```


# Client Handling Multiple Descriptors



- read() call on both stdin and socket will block until data is available.
- **Consider a case:**
  - If client is blocked in waiting for user to enter data, meanwhile TCP receives FIN from server.
    - Server is down. So sending request is meaningless.
- How to handle uncertainty of availability of data on descriptors?



Normal course of actions

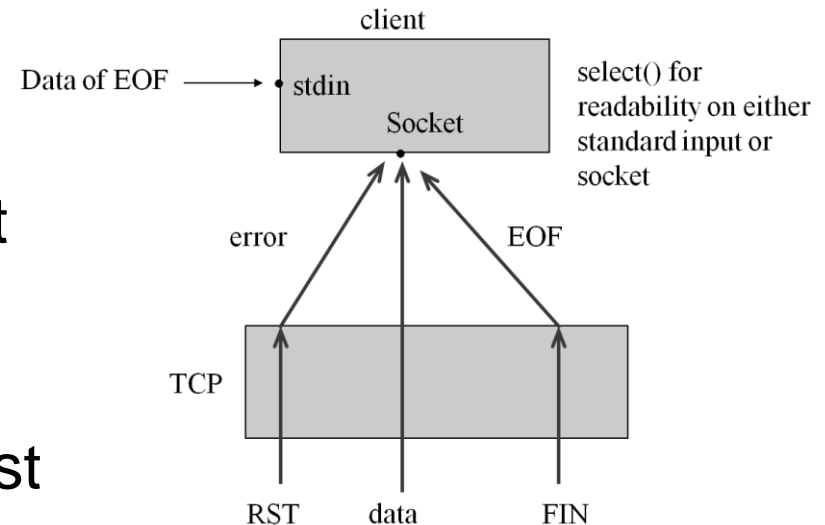


Unexpected Server shutdown

# read() on socket



- Peer TCP sends data, the socket becomes readable and *read* returns greater than 0.
- Peer TCP send a FIN(peer process terminates), the socket become readable and *read* returns 0(end-of-file)
- Peer TCP send a RST(peer host has crashed and rebooted), the socket become readable and returns -1 and *errno* contains the specific error code



# Client Handling Multiple Descriptors



- To avoid a situation where data has arrived from socket but client is unable to take note of it, use I/O Multiplexing.
- Client can wait on `select()`.
  - Add `stdin`, `socket` to `fd_set`.
  - Call `select ()` with `fd_set` for readability
  - When `select()` returns, find out which descriptor is ready with data.
  - Call `read()` on that `fd`.
- This will enable client to give timely response and avoid error situations.

# Client Handling Multiple Descriptors

innovate

achieve

lead

```
1 void str_cli(FILE *fp, int sockfd)
2 {
3     int maxfdp1;
4     fd_set rset;
5     char sendline[MAXLINE], recvline[MAXLINE];
6     FD_ZERO(&rset);
7     for ( ; ; ) {
8         FD_SET(fileno(fp), &rset);
9         FD_SET(sockfd, &rset);
10        maxfdp1 = max(fileno(fp), sockfd) + 1;
11        select(maxfdp1, &rset, NULL, NULL, NULL);
12        if (FD_ISSET(sockfd, &rset)) { /* socket is readable */
13            if (Readline(sockfd, recvline, MAXLINE) == 0)
14                err_quit("str_cli: server terminated prematurely");
15            Fputs(recvline, stdout);
16        }
17        if (FD_ISSET(fileno(fp), &rset)) { /* input is readable */
18            if (Fgets(sendline, MAXLINE, fp) == NULL)
19                return; /* all done */
20            Writen(sockfd, sendline, strlen(sendline));
21        }
22    } //for
23 } //str_cli
```

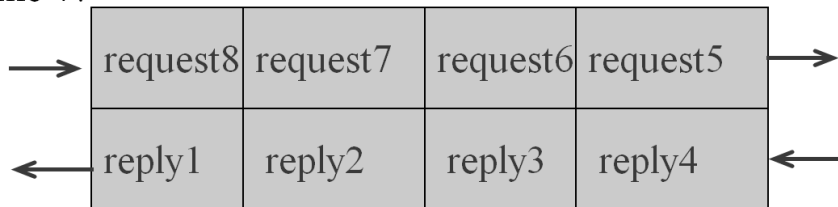


# Batch Mode Client

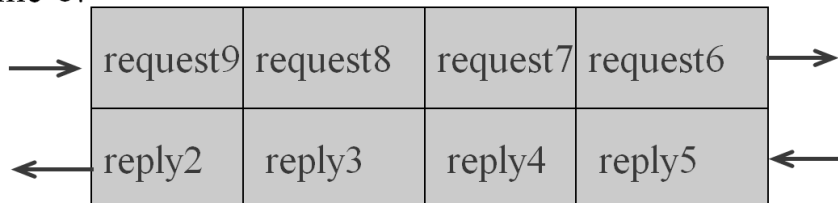


- Stop and Wait client
  - Send one request and wait for reply.
  - Usual in interactive mode.
- Batch Mode clients
  - Send requests without waiting for reply.
  - Better bandwidth utilization.

Time 7:



Time 8:



- Need for closing a socket partially:
  - We tell server that we have sent all requests by closing socket.
    - It will send FIN to server.
  - But in batch mode, by closing socket, we send FIN but we can't read replies which are yet to reach the client.
- `close()` vs `shutdown()` sys calls
  - closes the socket partially (either read end or write end) unlike `close()` sys call.
    - `close()` closes completely.
  - Irrespective of reference count it closes the socket.
    - `close()` will initiate FIN only if reference count for the fd reaches 0.

# shutdown() sys call



- Sometimes, it is useful to close one half of the connection, so that data can be transmitted in just one direction through the socket.

```
1  #include <sys/socket.h>
2  int shutdown(int sockfd , int how );
3  //Returns 0 on success, or -1 on error
```

- SHUT\_RD : read-half of the connection closed. Subsequent reads will return end-of-file (0).
  - SHUT\_RD can't be used meaningfully for TCP sockets.
- SHUT\_WR : write-half of the connection closed. Also called *socket half-close*. Buffered data will be sent followed by termination sequence.
  - Common use of shutdown()
  - Subsequent writes to the local socket yield the SIGPIPE signal and an EPIPE error.
- SHUT\_RDWR : both closed
  - Note that shutdown() doesn't close the file descriptor, even if how is specified as SHUT\_RDWR . To close the file descriptor, we must additionally call close().


# Batch Mode Client



- After user presses, Ctrl-D (EOF), close write half of the socket.
- Also set *stdineof* variable to 1.
  - This will help in inferring the FIN received from server as normal or abnormal termination.
  - In case of normal termination, we received all the replies.


EOF on *stdin*  
Set *stdineof*=1  
*shutdown*(fd, SHUT\_WR)

if *stdineof*=1  
Receiving FIN  
from server is  
Normal  
termination



Data on *stdin*

if *stdineof*=0  
Receiving FIN  
from server is  
abnormal  
termination





```
1 str_cli(FILE *fp, int sockfd)
2 {
3     int      maxfdp1, stdineof;
4     fd_set   rset;
5     stdineof = 0;
6     FD_ZERO(&rset);
7     for ( ; ; ) {
8         if (stdineof == 0)
9             FD_SET(fileno(fp), &rset);
10        FD_SET(sockfd, &rset);
11        maxfdp1 = max(fileno(fp), sockfd) + 1;
12        select(maxfdp1, &rset, NULL, NULL, NULL);
13        if (FD_ISSET(sockfd, &rset)) { /* socket is readable */
14            if ( (n = Read(sockfd, buf, MAXLINE)) == 0) {
15                if (stdineof == 1)
16                    return; /* normal termination */
17                else
18                    err_quit("str_cli: server terminated prematurely");
19            }
20            Write(fileno(stdout), buf, n);
21        }
22        if (FD_ISSET(fileno(fp), &rset)) { /* input is readable */
23            if ( (n = Read(fileno(fp), buf, MAXLINE)) == 0) {
24                stdineof = 1;
25                shutdown(sockfd, SHUT_WR); /* send FIN */
26                FD_CLR(fileno(fp), &rset);
27                continue;
28            }
29            Writen(sockfd, buf, n);
30        }
31    }
```

# TCP Server Using select()




- Single process server that uses select to handle any number of clients, instead of forking one child per client.
- Protocol: echo
- Two data structures:
  - Client array
    - Keeps list of client sockets connected currently
  - fd\_set *allset*
    - Keeps list of fds for checking against readability.

# TCP Server Using select()



- There are three fds: 0,1,2
- One more fd after creating listening socket.

Before first client has established a connection




# TCP Server Using select()



- When a new client is accepted through accept()
  - A connected socket is added

After first client connection is established






# TCP Server Using select()



- When second client is accepted

After second client connection is established

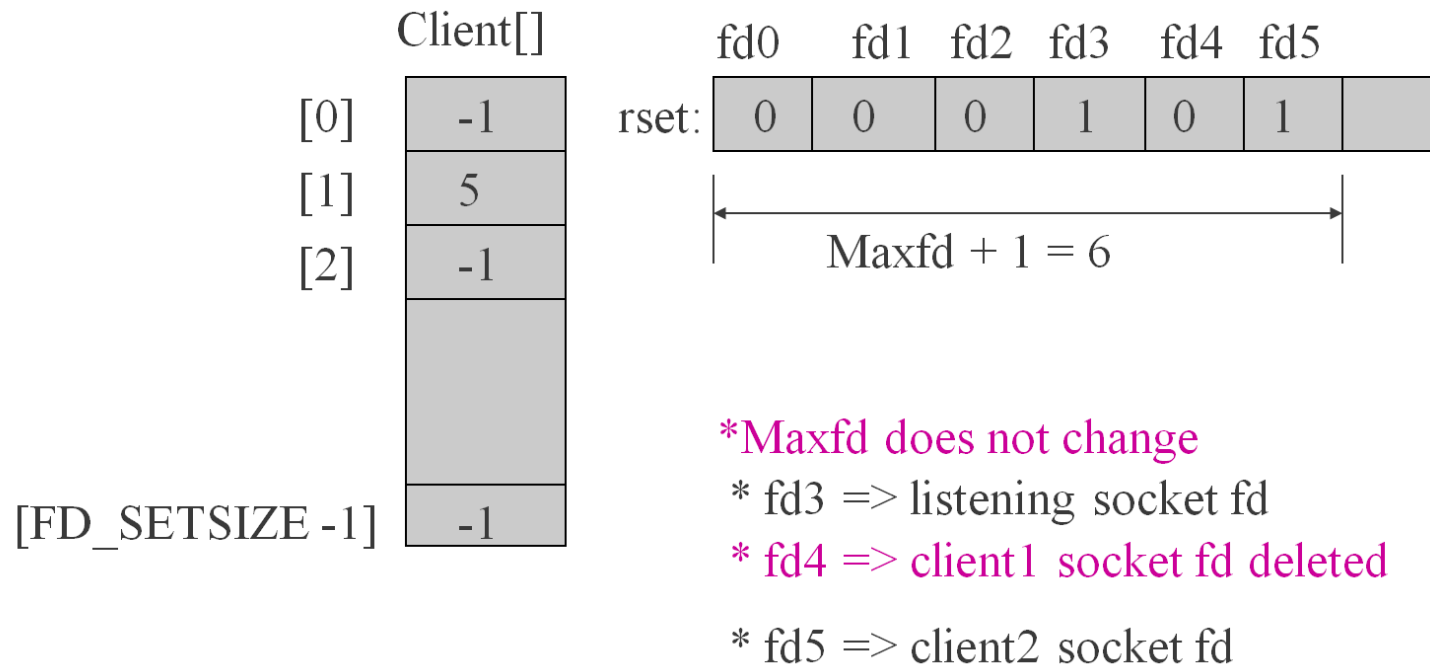


# TCP Server Using select()



- When the first client terminates connection
  - This is known when read() returns zero.

After first client terminates its connection



# TCP Server Using select()



- Create a passive socket.

```
1  int main(int argc, char **argv)
2  {
3      int          i, maxi, maxfd, listenfd, connfd, sockfd;
4      int          nready, client[FD_SETSIZE];
5      fd_set       rset, allset;
6      struct sockaddr_in cliaddr, servaddr;
7      listenfd = socket(AF_INET, SOCK_STREAM, 0);
8      bzero(&servaddr, sizeof(servaddr));
9      servaddr.sin_family      = AF_INET;
10     servaddr.sin_addr.s_addr = htonl(INADDR_ANY);
11     servaddr.sin_port        = htons(SERV_PORT);
12     bind(listenfd, (SA *) &servaddr, sizeof(servaddr));
13     listen(listenfd, LISTENQ);
```

# TCP Server Using select()



- Handling when listening socket is readable

```
1  maxfd = listenfd;           /* initialize */
2  maxi = -1;                  /* index into client[] array */
3  for (i = 0; i < FD_SETSIZE; i++)
4  client[i] = -1;             /* -1 indicates available entry */
5  FD_ZERO(&allset);
6  FD_SET(listenfd, &allset);
7  for ( ; ; ) {
8      rset = allset;          /* structure assignment */
9      nready = select(maxfd+1, &rset, NULL, NULL, NULL);

10     if (FD_ISSET(listenfd, &rset)) { /* new client connection */
11         cliilen = sizeof(cliaddr);
12         connfd = accept(listenfd, (SA *) &cliaddr, &cliilen);
13         for (i = 0; i < FD_SETSIZE; i++)
14             if (client[i] < 0) {
15                 client[i] = connfd;    /* save descriptor */
16                 break;
17             }
18         if (i == FD_SETSIZE) err_quit("too many clients");
19         FD_SET(connfd, &allset);      /* add new descriptor to set */
20         if (connfd > maxfd)
21             maxfd = connfd;           /* maxfd for select */
22         if (i > maxi)
23             maxi = i;                 /* max index in client[] array */
24         if (--nready <= 0)
25             continue;                /* no more readable descriptors */
26     }
```

# TCP Server Using select()



- When a connected socket is readable

```
1  for (i = 0; i <= maxi; i++) { /* check all clients for data */
2      if ((sockfd = client[i]) < 0)
3          continue;
4      if (FD_ISSET(sockfd, &rset)) {
5          if ( (n = Readline(sockfd, line, MAXLINE)) == 0) {
6              /*connection closed by client */
7              close(sockfd);
8              FD_CLR(sockfd, &allset);
9              client[i] = -1;
10         }
11         else
12             Writen(sockfd, line, n);
13         if (--nready <= 0)
14             break; /* no more readable descriptors */
15     }
16 }
```

- This code looks complicated when compared to fork-per-client model. But this design avoids overhead of fork().

# Denial-of-Service Attacks



- If malicious client connect to the server, send 1 byte of data (other than a newline), and then goes to sleep.
  - in readline(), server is blocked.
- Solution
  - use nonblocking I/O
  - have each client serviced by a separate thread of control (spawn a process or a thread to service each client)
  - place a timeout on the I/O operation

# pselect()

```
1 struct timespec{
2     time_t  tv_sec;    /*seconds*/
3     long    tv_nsec;   /* nanoseconds */
4 };
```



- pselect contains two changes from the normal select function:
  - pselect uses the timespec structure instead of the timeval structure.
  - Accepts signal mask.

```
1 #define _XOPEN_SOURCE 600
2 #include <sys/select.h>
3 int pselect(int  nfds , fd_set * readfds , fd_set * writefds ,
4     fd_set * exceptfds, struct timespec * timeout , const sigset_t * sigmask );
5 //Returns number of ready file descriptors, 0 on timeout, or -1 on error
```

```
2 ready = pselect(nfds, &readfds, &writefds, &exceptfds, timeout, &sigmask);
```

- This call is equivalent to

```
1 sigset_t origmask;
2 sigprocmask(SIG_SETMASK, &sigmask, &origmask);
3 ready = select(nfds, &readfds, &writefds, &exceptfds, timeout);
4 sigprocmask(SIG_SETMASK, &origmask, NULL);          /* Restore signal mask */
```

# Problems with select() and poll()



- The select() and poll() system calls are the portable, long-standing, and widely used methods of monitoring multiple file descriptors for readiness.
- Suffer from some problems
  - Kernel must check all the fds to check if they are ready.
  - Each time select() passes data structures which kernel modifies and returns.
  - Once select() returns, the program must inspect the data structure to see which fds are ready.
- Select() scales poorly with the increase of *fds*.
- Signal driven I/O or *epoll* (event poll) provide a scalable solution.





**BITS Pilani**  
Pilani Campus



# Non-blocking I/O on Sockets

- Input operations: read, recv, readv, recvfrom, recvmsg
  - Blocking operations
  - TCP: until a byte arrives.
  - UDP: until a datagram arrives.
  - With non-blocking socket, if no data, return with EWOULDBLOCK error.
- Output Operations: write, send, writev, sendto, sendmsg
  - Blocks if there is no room in socket send buffer.
  - TCP: until all the data is written.
  - UDP: no send buffer present.
  - With non-blocking socket, TCP write will write whatever it can and returns no. of bytes written. If no room at all, it returns with error EWOULDBLOCK.

# Socket Operations




- Accepting incoming connections: `accept`
  - Blocks if no incoming connection.
  - With non-blocking socket, it would return with an error.
- Initiating Connections: `connect`
  - Blocks until client TCP receives ACK.
  - With non-blocking socket, it returns *errno* EINPROGRESS, and continues to establish connection.

# Client Handling a socket, stdin, stdout



- A client usually deals with
  - Stdin
  - Stdout
  - Socket



# Client Handling a socket, stdin, stdout



- We looked at
  - Stop and wait client
  - Batch mode client - select with blocking I/O
    - Once select() returns, and if socket is readable, read() is called on socket.
    - readline() call gets blocked on socket till it gets required data.
    - During this time, other clients have to wait.
- Now we look at select with non-blocking I/O
  - In this, read() will be a non-blocking operation. It will read whatever data available on socket. It returns.
  - When this fd is readable next time, further data is read.
    - This requires that we track the number of bytes read and the pointers in the buffer.



```
1 str_cli(FILE *fp, int sockfd)
2 {
3     int      maxfdp1, stdineof;
4     fd_set   rset;
5     stdineof = 0;
6     FD_ZERO(&rset);
7     for ( ; ; ) {
8         if (stdineof == 0)
9             FD_SET(fileno(fp), &rset);
10        FD_SET(sockfd, &rset);
11        maxfdp1 = max(fileno(fp), sockfd) + 1;
12        select(maxfdp1, &rset, NULL, NULL, NULL);
13        if (FD_ISSET(sockfd, &rset)) { /* socket is readable */
14            if ( (n = Read(sockfd, buf, MAXLINE)) == 0) {
15                if (stdineof == 1)
16                    return; /* normal termination */
17                else
18                    err_quit("str_cli: server terminated prematurely");
19            }
20            Write(fileno(stdout), buf, n);
21        }
22        if (FD_ISSET(fileno(fp), &rset)) { /* input is readable */
23            if ( (n = Read(fileno(fp), buf, MAXLINE)) == 0) {
24                stdineof = 1;
25                shutdown(sockfd, SHUT_WR); /* send FIN */
26                FD_CLR(fileno(fp), &rset);
27                continue;
28            }
29            Writen(sockfd, buf, n);
30        }
31    }
```

# Select with Non-Blocking IO



- Non-blocking IO complicates buffer management.
  - We have to keep track of how much is read and how much is written.
- Two buffers:
  - *to*: reading from standard input and write to socket.
  - *from*: read from socket and write to stdout.

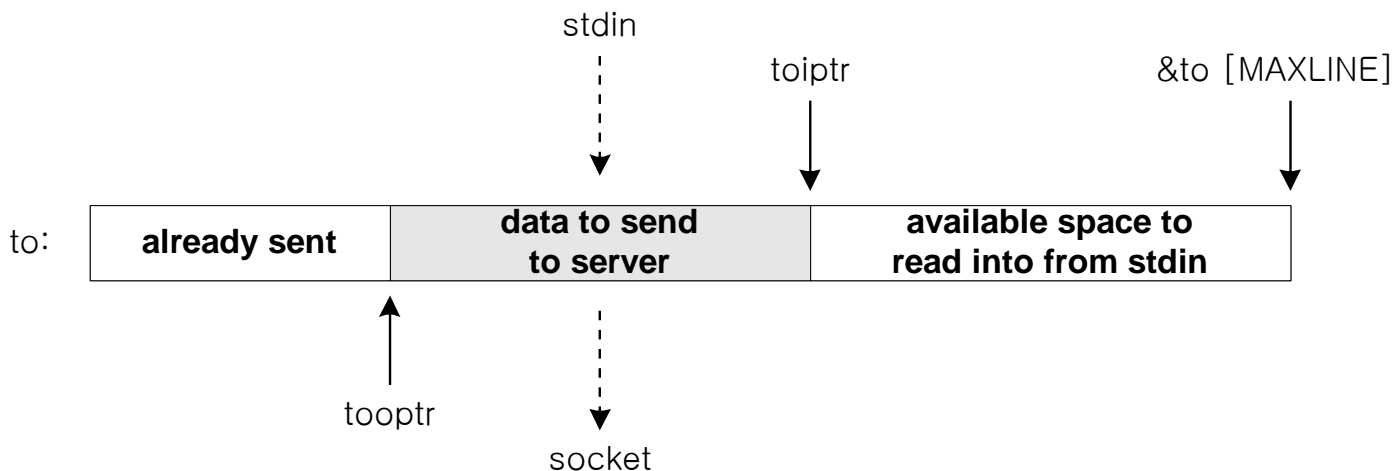


Figure 15.1 Buffer containing data from standard input going to the socket.

# Select with Non-Blocking IO

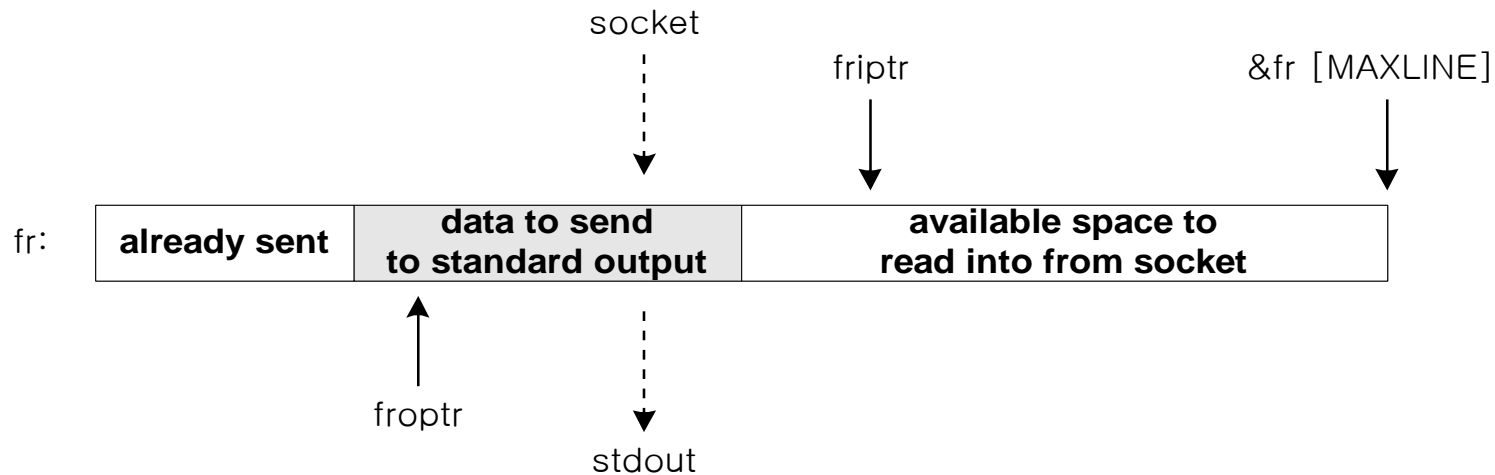


Figure 15.2 Buffer containing data from the socket going to standard output.

- *froptr*: points to the next byte to be sent to stdout.
- *friptr*: points to the next byte into which next byte can be read..



# Select with Non-Blocking IO



```
1 void str_cli(FILE *fp, int sockfd)
2 {
3     int      maxfdp1, val, stdineof;
4     ssize_t  n, nwritten;
5     fd_set   rset, wset;
6     char     to[MAXLINE], fr[MAXLINE];
7     char     *toiptr, *tooptr, *friptr, *froptr;
8     val = fcntl(sockfd, F_GETFL, 0);
9     fcntl(sockfd, F_SETFL, val | O_NONBLOCK);
10    val = Fcntl(STDIN_FILENO, F_GETFL, 0);
11    fcntl(STDIN_FILENO, F_SETFL, val | O_NONBLOCK);
12    val = Fcntl(STDOUT_FILENO, F_GETFL, 0);
13    fcntl(STDOUT_FILENO, F_SETFL, val | O_NONBLOCK);
14    toiptr = tooptr = to;        /* initialize buffer pointers */
15    friptr = froptr = fr;
16    stdineof = 0;
```

```
18    maxfdp1 = max(max(STDIN_FILENO, STDOUT_FILENO), sockfd) + 1;
19    for ( ; ; ) {
20        FD_ZERO(&rset);
21        FD_ZERO(&wset);
22        if (stdineof == 0 && toiptr < &to[MAXLINE])
23            FD_SET(STDIN_FILENO, &rset);    /* read from stdin */
24        if (friptr < &fr[MAXLINE])
25            FD_SET(sockfd, &rset);    /* read from socket */
26        if (tooptr != toiptr)
27            FD_SET(sockfd, &wset);    /* data to write to socket */
28        if (froptr != friptr)
29            FD_SET(STDOUT_FILENO, &wset);    /* data to write to stdout */
30        select(maxfdp1, &rset, &wset, NULL, NULL);
```

# reads from standard input



```
30 select(maxfdp1, &rset, &wset, NULL, NULL);
31 if (FD_ISSET(STDIN_FILENO, &rset)) {
32     if((n = read(STDIN_FILENO, toiptr, &to[MAXLINE] - toiptr)) < 0) {
33         if (errno != EWOULDBLOCK)
34             err_sys("read error on stdin");
35     } else if (n == 0) {
36         fprintf(stderr, "%s: EOF on stdin\n", gf_time());
37         stdineof = 1; /* all done with stdin */
38         if (tooptr == toiptr)
39             shutdown(sockfd, SHUT_WR); /* send FIN */
40     } else {
41         fprintf(stderr, "%s: read %d bytes from stdin\n", gf_time(),
42                 n);
43         toiptr += n; /* # just read */
44         FD_SET(sockfd, &wset); /* try and write to socket below */
45     }
46 }
```

Amt of space available in to buffer

If user has pressed Ctrl-D,  
set stdineof=1  
If no outstanding data on  
buffer, close the write end.

Increment to pointer  
set socket in wset for writability

# reads from socket

innovate

achieve

lead

Amt of space available in from buffer

```
47 if (FD_ISSET(sockfd, &rset)) {
48     if ( (n = read(sockfd, friptr, &fr[MAXLINE] - friptr)) < 0) {
49         if (errno != EWOULDBLOCK)
50             err_sys("read error on socket");
51     } else if (n == 0) {
52         fprintf(stderr, "%s: EOF on socket\n", gf_time());
53         if (stdineof)
54             return;      /* normal termination */
55         else
56             err_quit("str_cli: server terminated prematurely");
57     } else {
58         fprintf(stderr, "%s: read %d bytes from socket\n",
59                 gf_time(), n);
60         friptr += n;      /* # just read */
61         FD_SET(STDOUT_FILENO, &wset);    /* try and write below */
62     }
63 }
```

Increment friptr.  
Add stdout to wset to test for writability.

# writes to standard output

innovate

achieve

lead

No of bytes to write >0

```
65 ▾ if (FD_ISSET(STDOUT_FILENO, &wset) && ((n = friptr - froptr) > 0)) {
66 ▾     if ( (nwritten = write(STDOUT_FILENO, froptr, n)) < 0) {
67         if (errno != EWOULDBLOCK)
68             err_sys("write error to stdout");
69 ▾     } else {
70         fprintf(stderr, "%s: wrote %d bytes to stdout\n",
71             gf_time(), nwritten);
72         froptr += nwritten; /* # just written */
73         if (froptr == friptr)
74             froptr = friptr = fr; /* back to beginning of buffer */
75     }
76 }
```

If the write is successful, froptr is incremented by the number of bytes written

# Writes to socket

innovate

achieve

lead

No of bytes to write >0

```
78  if (FD_ISSET(sockfd, &wset) && ((n = toiptr - tooptr) > 0)) {
79      if ( (nwritten = write(sockfd, tooptr, n)) < 0) {
80          if (errno != EWOULDBLOCK)
81              err_sys("write error to socket");
82      } else {
83          fprintf(stderr, "%s: wrote %d bytes to socket\n",
84                  gf_time(), nwritten);
85          tooptr += nwritten; /* # just written */
86          if (tooptr == toiptr) {
87              toiptr = tooptr = to; /* back to beginning of buffer */
88              if (stdineof)
89                  shutdown(sockfd, SHUT_WR); /* send FIN */
90          }
91      }
92  }
93  }
94  }
```

If the write is successful, tooptr is incremented by the number of bytes written

if we encountered an EOF on standard input, the FIN can be sent to the server




# Client using Multiple Processes

# Client using Multiple Processes



- Whenever we find the need to use nonblocking I/O, it will usually be simpler to split the application into either processes (using fork) or threads.
  - Parent reads from stdin and writes to socket
  - Child reads from socket and writes to stdout.



- Normal termination occurs when the EOF on standard input is encountered. The parent reads this EOF and calls shutdown to send a FIN.
- If abnormal occurs, the child will read an EOF on the socket. If this happens, the child must tell the parent to stop copying from the standard input to the socket
  - the child sends a signal (e.g. SIGTERM) to the parent.

# Client using Multiple Processes

innovate

achieve

lead

```
1 void str_cli(FILE *fp, int sockfd)
2 {
3     pid_t    pid;
4     char      sendline[MAXLINE], recvline[MAXLINE];
5     if ( (pid = fork()) == 0) { /* child: server -> stdout */
6         while (Readline(sockfd, recvline, MAXLINE) > 0)
7             fputs(recvline, stdout);
8         kill(getppid(), SIGTERM); /* in case parent still running */
9         exit(0);
10    }
11    /* parent: stdin -> server */
12    while (fgets(sendline, MAXLINE, fp) != NULL)
13        Writen(sockfd, sendline, strlen(sendline));
14    shutdown(sockfd, SHUT_WR); /* EOF on stdin, send FIN */
15    pause();
16    return;
17 }
```



# Comparing Cleint Designs



- when copying 2,000 lines from a client to a server with an RTT of 175 ms:

Client	Time taken for sending and receiving
stop-and-wait	354.0 sec
select and blocking I/O	12.3 sec
nonblocking I/O	6.9 sec
fork	8.7 sec
threaded version	8.5 sec

- nonblocking I/O version is almost twice as fast as version using blocking I/O with select.
- Version using fork is slower than nonblocking I/O version.
- Nevertheless, given the complexity of the nonblocking I/O code versus the fork code, fork version is simple approach.



# Non-blocking Connect

# Nonblocking connect()



- TCP socket nonblocking connect
  - return: an error of EINPROGRESS
  - TCP three-way handshake continues
  - check the connection establishment using select
- There are three uses for a nonblocking connect.
  - We can overlap other processing with the three-way handshake.
  - We can establish multiple connections at the same time using this technique.
    - popular with Web browsers
  - Since we wait for the connection establishment to complete using select, we can specify time limit for select, allowing us to shorten the timeout for the connect.

# Nonblocking connect()



- Set the socket to non-blocking.
- Call `connect()`. It will return immediately with error `EINPROGRESS`.
- We use `select()` to check what has happened to `connect()`.
- If the descriptor is readable or writable, we call `getsockopt()` to fetch the socket's pending error (`SO_ERROR`). If the connection completed successfully, this value will be 0.

# Nonblocking connect()



```
1 int connect_nonb(int sockfd, const SA *saptr, socklen_t salen, int nsec)
2 {
3     int flags, n, error;
4     socklen_t len;
5     fd_set rset, wset;
6     struct timeval tval;
7     flags = fcntl(sockfd, F_GETFL, 0);
8     fcntl(sockfd, F_SETFL, flags | O_NONBLOCK);
9     error = 0;
10    if ( (n = connect(sockfd, saptr, salen)) < 0)
11        if (errno != EINPROGRESS)
12            return (-1);
13    /* Do whatever we want while the connect is taking place. */
14    if (n == 0)
15        goto done; /* connect completed immediately */
16    FD_ZERO(&rset);
17    FD_SET(sockfd, &rset);
18    wset = rset;
19    tval.tv_sec = nsec;
20    tval.tv_usec = 0;
```

# Nonblocking connect()

innovate

achieve

lead

```
22     if ( (n = select(sockfd + 1, &rset, &wset, NULL,
23         nsec ? &tval : NULL)) == 0) {
24         close(sockfd);          /* timeout */
25         errno = ETIMEDOUT;
26         return (-1);
27     }
28     if (FD_ISSET(sockfd, &rset) || FD_ISSET(sockfd, &wset)) {
29         len = sizeof(error);
30         if (getsockopt(sockfd, SOL_SOCKET, SO_ERROR, &error, &len) < 0)
31             return (-1);      /* Solaris pending error */
32     } else
33         err_quit("select error: sockfd not set");
34 done:
35     Fcntl(sockfd, F_SETFL, flags); /* restore file status flags */
36     if (error) {
37         close(sockfd);          /* just in case */
38         errno = error;
39         return (-1);
40     }
41     return (0);
42 }
```



**BITS Pilani**  
Pilani Campus



# Web client: Non-blocking Connect


# nonblocking connect: web client




- A real-world example of nonblocking connects started with Netscape Web Client
- The client establishes an HTTP connection with a Web server and fetches a home page.
- On that page are often numerous references to other Web pages.
- Instead of fetching these other pages serially, one at a time, the client can fetch more than one at the same time, using nonblocking connects.




# nonblocking connect: web client (cont.)



three connections  
done serially



three connections  
done in parallel;  
maximum of two  
connections at a time



three connections  
done in parallel;  
maximum of three  
connections at a time

# Non-blocking Connect



- This program will read up to 20 files from a Web server.
- We specify as command-line arguments
  - the maximum number of parallel connections,
  - the server's hostname, and
  - each of the filenames to fetch from the server.

```
2  bash$ web 3 www.foobar.com image1.gif image2.gif image3.gif image4.gif
3  image5.gif image6.gif image7.gif
```

- It means
  - three simultaneous connection
  - server's hostname
  - filename for the home page
  - the files to be read

```
1  #define MAXFILES      20
2  #define SERV          "80"          /* port number or service name */
3  struct file {
4      char    *f_name;                /* filename */
5      char    *f_host;                /* hostname or IPv4/IPv6 address */
6      int     f_fd;                   /* descriptor */
7      int     f_flags;                /* F_xxx below */
8  } file[MAXFILES];
9  #define F_CONNECTING    1           /* connect() in progress */
10 #define F_READING       2           /* connect() complete; now reading */
11 #define F_DONE          4           /* all done */
12 #define GET_CMD         "GET %s HTTP/1.0\r\n\r\n"
13 /* globals */
14 int     nconn, nfiles, nlefttoconn, nlefttoread, maxfd;
15 fd_set  rset, wset;
16 /* function prototypes */
17 void    home_page(const char *, const char *);
18 void    start_connect(struct file *);
19 void    write_get_cmd(struct file *);
```

- Each file has a state and fd.

# nonblock/web.c



```
1 main(int argc, char **argv)
2 {
3     int i, fd, n, maxnconn, flags, error;
4     char buf[MAXLINE];
5     fd_set rs, ws;
6     if (argc < 5)
7         err_quit("usage: web <#conns> <hostname> <homepage> <file1> ...");
8     maxnconn = atoi(argv[1]);
9     nfiles = min(argc - 4, MAXFILES);
10    for (i = 0; i < nfiles; i++) {
11        file[i].f_name = argv[i + 4];
12        file[i].f_host = argv[2];
13        file[i].f_flags = 0;
14    }
15    printf("nfiles = %d\n", nfiles);
16    home_page(argv[2], argv[3]);
17    FD_ZERO(&rset);
18    FD_ZERO(&wset);
19    maxfd = -1;
20    nlefttoread = nlefttoconn = nfiles;
21    nconn = 0;
```

- Process command-line arguments
- Read home page
- Initialize globals
  - Fd sets, nconn is current number of connections.

# nonblock/start\_connect.c



```
1 void start_connect(struct file *fptr)
2 {
3     int    fd, flags, n;
4     ai = Host_serv(fptr->f_host, SERV, 0, SOCK_STREAM);
5     fd = Socket(ai->ai_family, ai->ai_socktype, ai->ai_protocol);
6     fptr->f_fd = fd;
7     printf("start_connect for %s, fd %d\n", fptr->f_name, fd);
8     /* Set socket nonblocking */
9     flags = Fcntl(fd, F_GETFL, 0);
10    Fcntl(fd, F_SETFL, flags | O_NONBLOCK);
11    /* Initiate nonblocking connect to the server. */
12    if ( (n = connect(fd, ai->ai_addr, ai->ai_addrlen)) < 0) {
13        if (errno != EINPROGRESS)
14            err_sys("nonblocking connect error");
15        fptr->f_flags = F_CONNECTING;
16        FD_SET(fd, &rset);    /* select for reading and writing */
17        FD_SET(fd, &wset);
18        if (fd > maxfd)
19            maxfd = fd;
20    } else if (n >= 0)        /* connect is already done */
21        write_get_cmd(fptr); /* write() the GET command */
22 }
```

- Initiate nonblocking connect
- Handle connection complete
- If connect returns successfully, the connection is already complete and the function write\_get\_cmd ends a command to the server.

# nonblock/write\_get\_cmd.c



```
1  #include    "web.h"
2  void write_get_cmd(struct file *fptr)
3  {
4      int      n;
5      char     line[MAXLINE];
6      n = snprintf(line, sizeof(line), GET_CMD, fptr->f_name);
7      Writen(fptr->f_fd, line, n);
8      printf("wrote %d bytes for %s\n", n, fptr->f_name);
9      fptr->f_flags = F_READING; /* clears F_CONNECTING */
10     FD_SET(fptr->f_fd, &rset); /* will read server's reply */
11     if (fptr->f_fd > maxfd)
12         maxfd = fptr->f_fd;
13 }
```

- Build command and send it
- Set flags

# Main function: web.c



```
1 while (nlefttoread > 0) {
2     while (nconn < maxnconn && nlefttoconn > 0) {
3         /* find a file to read */
4         for (i = 0; i < nfiles; i++)
5             if (file[i].f_flags == 0)
6                 break;
7         if (i == nfiles)
8             err_quit("nlefttoconn = %d but nothing found", nlefttoconn);
9         start_connect(&file[i]);
10        nconn++;
11        nlefttoconn--;
12    }
```

- Initiate another connection, if possible

# Main function: web.c



```
13  rs = rset;
14  ws = wset;
15  n = select(maxfd + 1, &rs, &ws, NULL, NULL);
16  for (i = 0; i < nfiles; i++) {
17      flags = file[i].f_flags;
18      if (flags == 0 || flags & F_DONE)
19          continue;
20      fd = file[i].f_fd;
21      if (flags & F_CONNECTING &&
22          (FD_ISSET(fd, &rs) || FD_ISSET(fd, &ws))) {
23          n = sizeof(error);
24          if (getsockopt(fd, SOL_SOCKET, SO_ERROR, &error, &n) < 0 ||
25              error != 0) {
26              err_ret("nonblocking connect failed for %s",
27                      file[i].f_name);
28          }
29          /* connection established */
30          printf("connection established for %s\n", file[i].f_name);
31          FD_CLR(fd, &ws); /* no more writability test */
32          write_get_cmd(&file[i]); /* write() the GET command */

```

- select waits for either readability or writability.
  - Descriptors that have a nonblocking connect in progress will be enabled in both sets, while descriptors with a completed connection that are waiting for data from the server will be enabled in just the read set.



# Main function: web.c



```
29 ▾      /* connection established */
30      printf("connection established for %s\n", file[i].f_name);
31      FD_CLR(fd, &wset); /* no more writeability test */
32      write_get_cmd(&file[i]); /* write() the GET command */
33 ▾ } else if (flags & F_READING && FD_ISSET(fd, &rs)) {
34 ▾     if ( (n = Read(fd, buf, sizeof(buf))) == 0) {
35         printf("end-of-file on %s\n", file[i].f_name);
36         Close(fd);
37         file[i].f_flags = F_DONE; /* clears F_READING */
38         FD_CLR(fd, &rset);
39         nconn--;
40         nlefttoread--;
41 ▾     } else {
42         printf("read %d bytes from %s\n", n, file[i].f_name);
43     }
44 }
45 }
46 }
47 exit(0);
48 }
```

- If the F\_READING flag is set and the descriptor is ready for reading, we call read.

# Performance of Nonblocking Connect



- Table shows the clock time required to fetch a Web server's home page, followed by nine image files from that server.
  - The RTT to the server is about 150 ms.
  - The home page size was 4,017 bytes and the average size of the 9 image files was 1,621 bytes.
  - TCP's segment size was 512 bytes.
- Most of the improvement is obtained with three simultaneous connections.

# simultaneou s connec tions	Clock time (seconds), non blocking	Clock time(sec s) Threads
1	6.0	6.3
2	4.1	4.2
3	3.0	3.1
4	2.8	3.0
5	2.5	2.7
6	2.4	2.5
7	2.3	2.3
8	2.2	2.3
9	2.0	2.2

# Acknowledgements

---



# Q&A





**BITS Pilani**  
Pilani Campus



**Thank You**